

DP 200 - Implementing a Data Platform Solution

Lab 5 - Working with Relational Data Stores in the Cloud

Estimated Time: 75 minutes

Pre-requisites: It is assumed that the case study for this lab has already been read. It is assumed that the content and lab for module 1: Azure for the Data Engineer has also been completed

Lab files: The files for this lab are located in the *Allfiles\Labfiles\Starter\DP-200.5* folder.

Lab overview

The students will be able to provision an Azure SQL Database and Azure Synapse Analytics server and be able to issue queries against one of the instances that are created. They will be also be able to integrate a data warehouse with a number of other data platform technologies and use PolyBase to load data from one data source into Azure Synapse Analytics.

Lab objectives

After completing this lab, you will be able to:

1. Use Azure SQL Database
2. Describe Azure Synapse Analytics
3. Create and query Azure Synapse Analytics
4. Use PolyBase to load data into Azure Synapse Analytics

Scenario

You are the senior data engineer at AdventureWorks, and you are working with your team to transition a relational database system from an on-premises SQL Server to a Azure SQL Database located in Azure. You will begin by creating an instance of Azure SQL Database with the company's sample database. Your intention is to hand this instance of to a junior data engineer to perform some testing of departmental databases.

You will then provision Azure Synapse Analytics server and test that the provisioning of the server is successful by testing a sample database with a series of queries. You will then use PolyBase to load a dimension table from Azure Blob to test that the integration of this data platform technology with Azure Synapse Analytics.

At the end of this lab, you will have:

1. Used Azure SQL Database
2. Described Azure Synapse Analytics
3. Created and queried Azure Synapse Analytics
4. Used PolyBase to load data into Azure Synapse Analytics

IMPORTANT: As you go through this lab, make a note of any issue(s) that you have encountered in any provisioning or configuration tasks and log it in the table in the document located at *\Labfiles\DP-200-Issues-Doc.docx*. Document the Lab number, note the technology, Describe the issue, and what was the resolution. Save this document as you will refer back to it in a later module.

Exercise 1: Use Azure SQL Database

Estimated Time: 15 minutes

Individual exercise

The main task for this exercise are as follows:

1. Create and configure a SQL Database instance.

Task 1: Create and configure a SQL Database instance.

1. In the Azure portal, navigate to the + **Create a resource** blade.
2. In the New screen, click the **Search the Marketplace** text box, and type the word **SQL Database**. Click **SQL Database** in the list that appears.
3. In the **SQL Database** screen, click **Create**.
4. From the **Create SQL Database** screen, create an Azure SQL Database with the following settings:
 - In the Project details section, type in the following information
 - **Subscription**: the name of the subscription you are using in this lab
 - **Resource group**: awrgstudxx, where xx are your initials.
 - Click on the **Additional setting** tab, click **Sample** . The AdventureworksLT sample database is selected automatically.
 - Click the **Basics** tab once this has been done.
 - In the Database details section, type in the following information
 - Database name: type in **AdventureworksLT**
 - Server: Create a new server by clicking **Create new** with the following settings and click on **OK**:
 - **Server name**: sqlservicexx, where xx are your initials
 - **Server admin login**: xxsqladmin, where xx are your initials
 - **Password**: Pa55w.rd
 - **Confirm Password**: Pa55w.rd
 - **Location**: choose a **location** near to you.
 - click on **OK**

New server

Microsoft

Server name *

sqlserviceto ✓

.database.windows.net

Server admin login *

ctosqladmin ✓

Password *

***** ✓

Confirm password *

***** ✓

Location *


(Europe) West Europe ✓

- Leave the remaining settings to their defaults, and then click on **OK**

Home > New > SQL Database > Create SQL Database

Create SQL Database

Microsoft

 Changing Basic options may reset selections you have made. Review all options prior to creating the resource.

[Basics](#) [Networking](#) [Additional settings](#) [Tags](#) [Review + create](#)

Create a SQL database with your preferred configurations. Complete the Basics tab then go to Review + Create to provision with smart defaults, or visit each tab to customize. [Learn more](#)

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * ⓘ

Resource group * ⓘ
[Create new](#)

Database details

Enter required settings for this database, including picking a logical server and configuring the compute and storage resources

Database name *

Server * ⓘ
[Create new](#)

Want to use SQL elastic pool? * ⓘ ☐ Yes ☒ No

Compute + storage * ⓘ **General Purpose**
Gen5, 2 vCores, 32 GB storage
[Configure database](#)

5. In the **Create SQL Database** blade, click **Review + create**.

6. After the validation of the **Create SQL Database*** blade, click **Create**.

Note: The provision will takes approximately 4 minutes.

Result: After you completed this exercise, you have an Azure SQL Database instance

Exercise 2: Describe Azure Synapse Analytics

Estimated Time: 15 minutes

Individual exercise

The main tasks for this exercise are as follows:

1. Create and configure a Azure Synapse Analytics instance.
2. Configure the Server Firewall
3. Pause the warehouse database

Task 1: Create and configure a Azure Synapse Analytics instance.

1. In the Azure portal, click on the link **home** at the top left of the screen.
2. In the Azure portal, click + **Create a resource**.
3. In the New blade, navigate to the **Search the Marketplace** text box, and type the word **Synapse**. Click **Azure Synapse Analytics (formerly SQL DW)** in the list that appears.
4. In the **Azure Synapse Analytics (formerly SQL DW)** blade, click **Create**.
5. From the **SQL Data Warehouse** blade, create an Azure Synapse Analytics with the following settings:
 - In the **Project details** section, type in the following information
 - **Subscription:** the name of the subscription you are using in this lab

- **Resource group:** awrgstudxx, where xx are your initials.
- In **Additional setting** tab, under data source, click **Sample**.
- Click the **Basics** tab once this has been done.
- In the Database details section, type in the following information
 - **Database warehouse name:** Warehousexx, where xx are your initials.
 - **Server:** Create a new server by clicking **Create new** with the following settings and click on **OK**:
 - **Server name:** dwhservicexx, where xx are your initials
 - **Server admin login:** xxsqladmin, where xx are your initials
 - **Password:** Pa55w.rd
 - **Confirm Password:** Pa55w.rd
 - **Location:** choose a **location** near to you.
 - Select the checkbox to Allow Azure services to access server
 - click on **OK**

New server

Microsoft

Server name *

dwhservicecto

database.windows.net

Server admin login *

ctosqladmin

Password *

Confirm password *

Location *

(Europe) West Europe

☒ Allow Azure services to access server

- o Performance Level: Click **Select performance level** and select **Gen2 DW100C**.

Configure performance

Configure your performance level that best fits your data warehousing needs.

Gen2

Offers the highest performance and storage scalability options for intensive workloads.

Starting at 1.51 USD / hour

Gen1

Offers the lowest compute scale options for less demanding workloads.

Not available ⓘ

Starting at -- / hour

Learn more about [performance level](#)

Scale your system ⓘ

DW100c

1.51 USD / hour

100 cDWU

Learn more about [pricing](#)

Price of your system ⓘ

- Click **Apply**. the following configuration is shown.

Home > New > Azure Synapse Analytics (formerly SQL DW) > SQL Data Warehouse

SQL Data Warehouse

Microsoft

⚠ Changing basic options may reset selections you have made. Please review all options prior to creating the database.

🔑 Welcome to Azure Synapse Analytics (formerly known as Azure SQL Data Warehouse). [Learn more.](#)

[Basics](#) * [Additional settings](#) * [Tags](#) [Review + create](#)

Create a SQL data warehouse with your preferred configurations. Complete the Basics tab then go to Review + Create to provision with smart defaults, or visit each tab to customize. [Learn more](#) ⓘ

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * ⓘ ▼

Resource group * ⓘ ▼

[Create new](#)

Data warehouse details

Enter required settings for this data warehouse, including picking a logical server and configuring the performance level.

Data warehouse name * ✓

Server * ⓘ ▼

[Create new](#)

Performance level * ⓘ **Gen2**
DW100c
[Select performance level](#)

6. In the SQL Data Warehouse screen, click **Review + create**.

7. In the SQL Data Warehouse blade, click **Create**.

Note: The provision will takes approximately 7 minutes.

Task 2: Configure the Server Firewall

1. In the Azure portal, in the blade, click **Resource groups**, and then click **awrgstudxx**, and then click on **awdlstudxx**, where **xx** are your initials
2. Click on **dwhservicexx**, where **xx** are your initials.
3. In the **dwhservicexx** screen, click on **Firewalls and virtual networks**.
4. In the **dwhservicexx - Firewalls and virtual networks** screen, click on the option **+ Add client IP**, and then click on **Save**. On the success screen click **OK**.

Home > Resource groups > awrgstudcto > dwhservicecto - Firewalls and virtual networks

dwhservicecto - Firewalls and virtual networks

SQL server

Search (Ctrl+/) Save Discard **+ Add client IP**

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Settings

Quick start

Failover groups

Manage Backups

Active Directory admin

SQL databases

SQL elastic pools

Deleted databases

Import/Export history

DTU quota

Connections from the IPs specified below provides access to all the databases in dwhservicecto.

Allow Azure services and resources to access this server
ON OFF

Client IP address 31.51.197.202

| Rule name | Start IP | End IP |
|---------------------------|---------------|---------------|
| ClientIPAddress_2019-12-1 | 31.51.197.202 | 31.51.197.202 |

Connections from the VNET/Subnet specified below provides access to all databases in dwhservicecto.

Virtual networks **+ Add existing virtual network** **+ Create new virtual network**

| Rule name | Virtual network | Subnet | Address Range | Endpoint status | Resource group | Subscription |
|--------------------------------|-----------------|--------|---------------|-----------------|----------------|--------------|
| No vnet rules for this server. | | | | | | |

Note: You will receive a message stating that the the server firewall rules have been successfully updated

5. Close down the Firewalls and virtual networks screen.

Result: After you completed this exercise, you have created an Azure Synapse Analytics instance and configures the server firewall to enable connections against it.

Task 3: Pause the Warehousexx database

1. Click on **Warehousexx**, where **xx** are your initials.
2. In the **Warehousexx (dwhservicexx/Warehousexx)** screen, click on **Pause**.
3. In the Pause Warehousexx screen, click **Yes**

Exercise 3: Creating an Azure Synapse Analytics database and tables

Estimated Time: 25 minutes

Individual exercise

The main tasks for this exercise are as follows:

1. Install SQL Server Management Studio and connect to a data warehouse instance.
2. Create a SQL Data Warehouse database
3. Create SQL Data Warehouse tables

Note: If you are not familiar with Transact-SQL, statements are available for the following labs in the following location
Allfiles\Labfiles\Starter\DP-200.5\SQL DW Files

Task 1: Install SQL Server Management Studio and connect to a SQL Data Warehouse instance.

1. In the Azure Portal, in the **dwhservicexx - Firewalls and virtual networks**, in the blade, click on **Properties**
2. Copy the "**Server name**" and paste it into Notepad.
3. Download [SQL Server Management Studio](#) and install onto your machine
4. On the windows desktop, click on the **Start**, and type "**SQL Server**" and then click on **Microsoft SQL Server Management Studio 17**
5. In the **Connect to Server** dialog box, fill in the following details
 - Server Name: **dwhservicexx.database.windows.net**
 - Authentication: **SQL Server Authentication**
 - Username: **xxsqladmin**
 - Password: **Pa55w.rd**
6. In the **Connect to Server** dialog box, click **Connect**

Task 2: Create a SQL Data Warehouse database.

1. In **SQL Server Management Studio**, in Object Explorer, right click **dwhservicexx.database.windows.net** and click on **New Query**.
2. In the query window, create a DataWarehouse database named **DWDB**, with a service objective of DW100 and a maximum size of 1024GB.

```
CREATE DATABASE DWDB COLLATE SQL_Latin1_General_CP1_CI_AS
(
    EDITION = 'DataWarehouse'
,    SERVICE_OBJECTIVE = 'DW100C'
,    MAXSIZE = 1024 GB
);
```

Note: The creation of the database takes approximately 2 minutes.

Task 3: Create SQL Data Warehouse tables.

1. In **SQL Server Management Studio**, in Object Explorer, right click **dwhservicexx.database.windows.net** and click on **New Query**.
2. In **SQL Server Management Studio**, in SQL Editor toolbar, in **Available Databases**, click on **DWDB**.

Note: If you are unfamiliar with Transact-SQL, there is a script in the Allfiles\Solution\DP-200.5\folder named **Exercise3 Task3Step2 script.sql**. It contains the bulk of the code required to create the tables, but you do have to complete the code by selecting the distribution type to use for each table

3. Create a table named **dbo.Users** with a **clustered columnstore** index with a distribution of **replicate** with the following columns:

| column name | data type | Nullability |
|-------------|---------------|-------------|
| userId | int | NULL |
| City | nvarchar(100) | NULL |
| Region | nvarchar(100) | NULL |
| Country | nvarchar(100) | NULL |

4. In **SQL Server Management Studio**, click on **Execute**.
5. In **SQL Server Management Studio**, in Object Explorer, right click **dwhservicexx.database.windows.net** and click on **New Query**.
6. In **SQL Server Management Studio**, in SQL Editor toolbar, in **Available Databases**, click on **DWDB**.
7. Create a table named **dbo.Products** with a **clustered columnstore** index with a distribution of **round robin** with the following columns:

| column name | data type | Nullability |
|--------------------|---------------|-------------|
| ProductId | int | NULL |
| EnglishProductName | nvarchar(100) | NULL |
| Color | nvarchar(100) | NULL |
| StandardCost | int | NULL |
| ListPrice | int | NULL |
| Size | nvarchar(100) | NULL |
| Weight | int | NULL |
| DaysToManufacture | int | NULL |
| Class | nvarchar(100) | NULL |
| Style | nvarchar(100) | NULL |

8. In **SQL Server Management Studio**, click on **Execute**.
9. In **SQL Server Management Studio**, in Object Explorer, right click **dwhservicexx.database.windows.net** and click on **New Query**.
10. In **SQL Server Management Studio**, in SQL Editor toolbar, in **Available Databases**, click on **DWDB**.
11. Create a table named **dbo.FactSales** with a **clustered columnstore** index with a distribution of **Hash** on the **SalesUnit** with the following columns:

| column name | data type | Nullability |
|-------------|-----------|-------------|
| DateId | int | NULL |
| ProductId | int | NULL |
| UserId | int | NULL |

| column name | data type | Nullability |
|------------------|-----------|-------------|
| UserPreferenceId | int | NULL |
| SalesUnit | int | NULL |

12. In **SQL Server Management Studio**, click on **Execute**.

Result: After you completed this exercise, you have installed SQL Server Management Studio to create a data warehouse named DWDB and three tables named Users, Products and FactSales.

Exercise 4: Using PolyBase to Load Data into Azure Synapse Analytics

Estimated Time: 10 minutes

Individual exercise

The main tasks for this exercise are as follows:

1. Collect Data Lake Storage container and key details
2. Create a dbo.Dates table using PolyBase from Azure Data Lake Storage

Task 1: Collect Azure Blob account name and key details

1. In the Azure portal, click on **Resource groups** and then click on **awrgstudxx**, and then click on **awdlsstudxx** where xx are the initials of your name.
2. In the **awdlsstudxx** screen, click **Access keys**. Click on the icon next to the **Storage account name** and paste it into Notepad.
3. In the **awdlsstudxx - Access keys** screen, under **key1**, Click on the icon next to the **Key** and paste it into Notepad.

Task 2: Create a dbo.Dates table using PolyBase from Azure Blob

1. In **SQL Server Management Studio**, in Object Explorer, right click **dwhservicexx.database.windows.net** and click on **New Query**.
2. In **SQL Server Management Studio**, in SQL Editor toolbar, in **Available Databases**, click on **DWDB**.
3. Create a **master key** against the **DWDB** database. In the query editor, type in the following code:

```
CREATE MASTER KEY;
```

4. Create a database scoped credential named **AzureStorageCredential** with the following details, by typing in the following code:
 - IDENTITY: **MOCID**
 - SECRET: **The access key of your storage account**

```
CREATE DATABASE SCOPED CREDENTIAL AzureStorageCredential
WITH
IDENTITY = 'MOCID',
SECRET = 'Your storage account key'
```

```
;
```

1. In **SQL Server Management Studio**, highlight both statements and then click on **Execute**.
2. In **SQL Server Management Studio**, in the Query window, type in code that will create an external data source named **AzureStorage** for the Blob storage account and data container created in with a type of **HADOOP** that makes use of the ****AzureStorageCredential**. Note that you should replace **awdlsstudxx** in the location key with your storage account with your initials

```
CREATE EXTERNAL DATA SOURCE AzureStorage
WITH (
  TYPE = HADOOP,
  LOCATION = 'abfs://data@awdlsstudxx.dfs.core.windows.net',
```



```
CREDENTIAL = AzureStorageCredential
);
```

3. In **SQL Server Management Studio**, in the Query window, type in code that will create an external file format named **TextFile** with a format type of **DelimitedText** and a field terminator of **comma**.

```
CREATE EXTERNAL FILE FORMAT TextFile
WITH (
    FORMAT_TYPE = DelimitedText,
    FORMAT_OPTIONS (FIELD_TERMINATOR = ',')
);
```

4. In **SQL Server Management Studio**, highlight the statement and then click on **Execute**.

5. In **SQL Server Management Studio**, in the Query window, type in code that will create an external table named **dbo.DimDate2External** with the **location** as the root file, the Data source as **AzureStorage**, the File_format of **TextFile** with the following columns:

| column name | data type | Nullability |
|---------------|----------------|-------------|
| Date | datetime2(3) | NULL |
| DateKey | decimal(38, 0) | NULL |
| MonthKey | decimal(38, 0) | NULL |
| Month | nvarchar(100) | NULL |
| Quarter | nvarchar(100) | NULL |
| Year | decimal(38, 0) | NULL |
| Year-Quarter | nvarchar(100) | NULL |
| Year-Month | nvarchar(100) | NULL |
| Year-MonthKey | nvarchar(100) | NULL |
| WeekDayKey | decimal(38, 0) | NULL |
| WeekDay | nvarchar(100) | NULL |
| Day Of Month | decimal(38, 0) | NULL |

```
CREATE EXTERNAL TABLE dbo.DimDate2External (
    [Date] datetime2(3) NULL,
    [DateKey] decimal(38, 0) NULL,
    [MonthKey] decimal(38, 0) NULL,
    [Month] nvarchar(100) NULL,
    [Quarter] nvarchar(100) NULL,
    [Year] decimal(38, 0) NULL,
    [Year-Quarter] nvarchar(100) NULL,
    [Year-Month] nvarchar(100) NULL,
    [Year-MonthKey] nvarchar(100) NULL,
    [WeekDayKey] decimal(38, 0) NULL,
    [WeekDay] nvarchar(100) NULL,
    [Day Of Month] decimal(38, 0) NULL
)
WITH (
    LOCATION='/DimDate2.txt',
    DATA_SOURCE=AzureStorage,
    FILE_FORMAT=TextFile
);
```

6. In **SQL Server Management Studio**, highlight the statement and then click on **Execute**.

7. Test that the table is created by running a select statement against it

```
SELECT * FROM dbo.DimDate2External;
```

8. In **SQL Server Management Studio**, in the Query window, type in a **CTAS** statement that creates a table named **dbo.Dates** with a **columnstore** index and a **distribution of round robin** that loads data from the **dbo.DimDate2External** table.

```
CREATE TABLE dbo.Dates
WITH
(
    CLUSTERED COLUMNSTORE INDEX,
    DISTRIBUTION = ROUND_ROBIN
)
AS
SELECT * FROM [dbo].[DimDate2External];
```

9. In **SQL Server Management Studio**, highlight the statement and then click on **Execute**.

10. In **SQL Server Management Studio**, in the Query window, type in a query that creates statistics on the **DateKey**, **Quarter** and **Month** column.

```
CREATE STATISTICS [DateKey] ON [Dates] ([DateKey]);
CREATE STATISTICS [Quarter] ON [Dates] ([Quarter]);
CREATE STATISTICS [Month] ON [Dates] ([Month]);
```

11. Test that the table is created by running a select statement against it

```
SELECT * FROM dbo.Dates;
```